

**Udacity AI Nanodegree**

**Project 3 – Adversarial Game Playing Agent**

Stephen Blystone

## Introduction

I chose the 3<sup>rd</sup> option for this project: Build an agent using advanced search techniques.

I wrote code for several search techniques including Negascout, Principal Variation Search (PVS), an iterative version of PVS, PVS with Zero Window Search, and Monte Carlo Tree Search.

I initially tried Negascout and PVS but due to poor performance I continued trying different search techniques before settling on Monte Carlo Tree Search (MCTS).

While I tested playing against a random agent and my own agent (self), for all the test results below I played against the MiniMax Agent.

## Performance Baseline

I wrote an Alpha Beta with Iterative Deepening search agent to use as my baseline against my MCTS agent. I configured a depth limit of 5 for the iterative deepening.

## Fairness

“Fair” matches are defined as repeating every game that is played but the agents switch initiative and use their opponent’s opening move. This is intended to balance out the advantage of picking “perfect openings”.

Prior to running any games my hypothesis was that running “fair” games would improve the win percentage due to some opening positions possibly being better than others. If the opponent started in a great opening position during a game, then they could have an unfair advantage and result in a loss for me. But by introducing “fairness” then my agent would have a chance to be in the same opening position, which could result in a win for me.

I tried running my agent and the baseline with “fair” and “unfair” games and found varied results. When I had an optimal value of  $c$  (0.25 or 0.5) then “fair” provided an advantage. When I had an unoptimal value of  $c$  (0, 0.75, 1, 100) then “fair” and “unfair” didn’t really make a difference.

## Search Time

I tried increasing the search time (the maximum allowed time for my Agent’s `get_action()` method to respond) and found that the results were generally improved. Some games didn’t appear to be impacted by the increased search time.

## MCTS Exploration vs Exploitation

Part of the power of the MCTS algorithm comes from the ability to tune the amount of exploration vs exploitation in the Best Child method. Exploration is when the agent will randomly explore in hopes of finding a good move. Exploitation is using the existing knowledge to make moves with known outcomes.

The equation is “exploitation + c \* exploration”, where c is the tunable parameter. I tested values of 0, 0.25, 0.5, 0.75, 1, and 100. When c is 0 there should be no exploration, when c is 1 then exploration and exploitation have equal importance, and when c is 100 then the exploration component should dominate over the exploitation.

I found the best values occurred with c = 0.25 and c = 0.5 indicating that a balance of exploration and exploitation is critical for success. If I were to extend this code further, I might explore using a variable parameter that starts off allowing more exploration and reduces the amount later in the game.

## Required Questions

- 1) How much performance difference does your agent show compared to the baseline?

Algorithm	# Rounds	Tunable Parameter for Exploration	Time Limit	Fairness?	% Wins
MCTS	100	0.5	25000	Fair	74.80%
Alpha Beta Iterative Deepening	100	N/A	25000	Fair	66.20%

Figure 1 - % Wins MCTS vs Baseline

My MCTS Agent’s best run was with the exploration tunable parameter set to 0.5, timer set to 25000ms, and running “fair” resulting in 74.8% wins. When using the baseline Alpha-Beta with Iterative Deepening with the same parameters it won 66.2%. My agent had an 8.6% higher win percentage over the baseline.

I modified my code to display the number of nodes searched for both the baseline and MCTS agent. The MCTS generally simulated more nodes than the baseline Alpha Beta agent was able to search. But looking at the number of nodes selected and expanded (which equals the number of nodes seen in backpropagation), there are significantly fewer nodes than the baseline. See appendix for counts.

- 2) Why do you think the technique you chose was more (or less) effective than the baseline?

I think my MCTS agent was more effective than the baseline because of the exploration/exploitation component of the algorithm. While the baseline Alpha-Beta with Iterative Deepening was limited to a certain depth, the MCTS agent was able to go deeper and simulate all the way to terminal state. As we saw in lecture, predicting a score at a fixed depth can vary as you go deeper one level. By being able to simulate further down the tree MCTS was able to make better move selections.

## Appendix – Full Data Tables

### Comparing % Wins for Various Parameter Settings

Algorithm	# Rounds	Tunable Parameter for Exploration	Time Limit	Fairness?	% Wins
MCTS	100	0	Default	Not Fair	63.00%
MCTS	100	0	Default	Fair	62.50%
MCTS	100	0	25000	Not Fair	65.00%
MCTS	100	0	25000	Fair	68.00%
MCTS	100	0.25	Default	Not Fair	67.00%
MCTS	100	0.25	Default	Fair	73.20%
MCTS	100	0.25	25000	Not Fair	67.00%
MCTS	100	0.25	25000	Fair	71.80%
MCTS	100	0.5	Default	Not Fair	68.50%
MCTS	100	0.5	Default	Fair	67.00%
MCTS	100	0.5	25000	Not Fair	70.50%
MCTS	100	0.5	25000	Fair	74.80%
MCTS	100	0.75	Default	Not Fair	63.00%
MCTS	100	0.75	Default	Fair	61.80%
MCTS	100	0.75	25000	Not Fair	61.00%
MCTS	100	0.75	25000	Fair	66.20%
MCTS	100	1	Default	Not Fair	60.00%
MCTS	100	1	Default	Fair	63.50%
MCTS	100	1	25000	Not Fair	66.50%
MCTS	100	1	25000	Fair	62.20%
MCTS	100	100	Default	Not Fair	17.50%
MCTS	100	100	Default	Fair	14.00%
MCTS	100	100	25000	Not Fair	18.50%
MCTS	100	100	25000	Fair	13.20%
Alpha Beta Iterative Deepening	100	N/A	Default	Not Fair	Not able to get results; Timeouts occurred and run_match.py stopped playing games.
Alpha Beta Iterative Deepening	100	N/A	Default	Fair	Not able to get results; Timeouts occurred and run_match.py stopped playing games.
Alpha Beta Iterative Deepening	100	N/A	25000	Not Fair	66.50%
Alpha Beta Iterative Deepening	100	N/A	25000	Fair	66.20%

### Comparing Number Nodes Searched

Game Move	Monte Carlo Tree Search (MCTS)		Alpha Beta
	MCTS Selected & Expanded	MCTS Simulated	Nodes Searched
1	416	4242	2232
2	397	4541	4865
3	531	3935	2398
4	393	3842	2935
5	428	3463	786
6	368	3384	5310
7	461	3043	1636
8	409	2642	785
9	428	2377	3855
10	416	2488	3008
11	431	2095	1221
12	613	2183	1732
13	573	2092	2373
14	509	2011	3713
15	467	1969	2179
16	528	2005	1107
17	554	1727	875
18	434	1751	1678
19	486	1486	554
20	541	1497	559
21	502	1482	507
22	440	1038	135
23	612	1020	211
24	540	908	456
25	523	840	396
26	885	679	275
27	915	357	197
28	782	175	263
29	574	56	161
30	406	8	110
31			226
32			112
33			61
34			24